

Raynal L, Marin J-M, Pudlo P, Ribatet M, Robert CP, Estoup A. 2017. **ABC random forests for Bayesian parameter inference.** *arXiv* 1605.05537v4, https://arxiv.org/pdf/1605.05537

A Preprint reviewed and recommended by *Peer Community In Evolutionary Biology*:

http://dx.doi.org/10.24072/pci.evolbiol.100036

ABC random forests for Bayesian parameter inference

Louis Raynal¹, Jean-Michel Marin*^{1,6}, Pierre Pudlo^{2,6}, Mathieu Ribatet¹, Christian P. Robert^{3,4}, and Arnaud Estoup^{5,6}

¹IMAG, Univ. Montpellier, CNRS, Montpellier, France
 ²Institut de Mathématiques de Marseille, Aix-Marseille Université, France
 ³Université Paris Dauphine, PSL Research University, Paris, France
 ⁴Departement of Statistics, University of Warwick, Coventry, UK
 ⁵CBGP, INRA, CIRAD, IRD, Montpellier SupAgro, Univ. Montpellier, Montpellier, France
 ⁶IBC, Univ. Montpellier, CNRS, Montpellier, France

Abstract

This preprint has been reviewed and recommended by Peer Community In Evolutionary Biology (http://dx.doi.org/10.24072/pci.evolbiol.100036). Approximate Bayesian computation (ABC) has grown into a standard methodology that manages Bayesian inference for models associated with intractable likelihood functions. Most ABC implementations require the preliminary selection of a vector of informative statistics summarizing raw data. Furthermore, in almost all existing implementations, the tolerance level that separates acceptance from rejection of simulated parameter values needs to be calibrated. We propose to conduct likelihood-free Bayesian inferences about parameters with no prior selection of the relevant components of the summary statistics and bypassing the derivation of the associated tolerance level. The approach relies on the random forest methodology of Breiman (2001) applied in a (non parametric) regression setting. We advocate the derivation of a new random forest for each component of the parameter vector of interest. When compared with earlier ABC solutions, this method offers significant gains in terms of robustness to the choice of the summary statistics, does not depend on any type of tolerance level, and is a good trade-off in term of quality of point estimator precision and credible interval estimations for a given computing time. We illustrate the performance of our methodological proposal and compare it with earlier ABC methods on a Normal toy example and a population genetics example dealing with human population evolution. All methods designed here have been incorporated in the R package abcrf (version 1.7) available on CRAN.

Keywords: Approximate Bayesian computation, Bayesian inference, likelihood-free methods, parameter inference, random forests

^{*}jean-michel.marin@umontpellier.fr

1 Introduction

As statistical models and data structures get increasingly complex, managing the likelihood function becomes a more and more frequent issue. We now face many realistic fully parametric situations where the likelihood function cannot be computed in a reasonable time or simply is unavailable. As a result, while the corresponding parametric model is well-defined, with unknown parameter θ , standard solutions based on the density function $f(y \mid \theta)$ like Bayesian or maximum likelihood analyses are prohibitive to implement. To bypass this hurdle, the last decade witnessed different inferential strategies, among which composite likelihoods (Lindsay, 1988; Varin et al., 2011), indirect inference (Gourieroux et al., 1993), GMMs (Chernozhukov and Hong, 2003), and likelihood-free methods such as approximate Bayesian computation (ABC, Beaumont et al., 2002; Csilléry et al., 2010; Marin et al., 2012), became popular options. We focus here on improving the latter solution.

Since their introduction in population genetics (Tavaré et al., 1997; Pritchard et al., 1999; Beaumont et al., 2002), ABC methods have been used in an ever increasing range of applications, corresponding to different types of complex models in diverse scientific fields (see, e.g., Beaumont, 2008; Toni et al., 2009; Beaumont, 2010; Csilléry et al., 2010; Theunert et al., 2012; Chan et al., 2014; Arenas et al., 2015; Sisson et al., 2017).

Posterior distributions are the cornerstone of any Bayesian analysis as they constitute both a sufficient summary of the data and a means to deliver all aspects of inference, from point estimators to predictions and uncertainty quantification. However, it is rather common that practitioners and users of Bayesian inference are not directly interested in the posterior distribution *per se*, but rather in some summary aspects, like posterior mean, posterior variance or posterior quantiles, since these are easier to interpret and report. With this motivation, we consider a version of ABC focussing on the approximation of unidimensional transforms of interest like the above, instead of resorting to the classical ABC approach that aims at approximating the entire posterior distribution and then handling it as in regular Bayesian inference. The approach we study here is based on random forests (RF, Breiman, 2001), which produces non-parametric regressions on an arbitrary set of potential regressors. We recall that the calibration side of RF (i.e. the choice of the RF parameters: typically the number of trees and the number of summary statistics sampled at each node) was successfully exploited in Pudlo et al. (2016) for conducting ABC model choice.

After exposing the ABC and RF principles, we explain how to fuse both methodologies towards Bayesian inference about parameters of interest. We then illustrate the performance of our proposal and compare it with earlier ABC methods on a Normal toy example and a population genetics example dealing with human population evolution.

2 Methods

Let

$$\{f(y \mid \theta) : y \in \mathcal{Y}, \theta \in \Theta\}, \ \mathcal{Y} \subseteq \mathbb{R}^n, \ \Theta \subseteq \mathbb{R}^p \qquad p, n \ge 1$$

be a parametric statistical model and $\pi(\theta)$ be a prior distribution on the parameter θ . Given an observation (or sample) y issued from this model, Bayesian parameter inference is based on the posterior distribution $\pi(\theta \mid y) \propto \pi(\theta) f(y \mid \theta)$. The computational difficulty addressed by ABC techniques is that a numerical evaluation of the density (a.k.a., likelihood) $f(y \mid \theta)$ is impossible or at least very costly, hence preventing the derivation of the posterior $\pi(\theta \mid y)$, even by techniques like MCMC (Marin and Robert, 2014).

2.1 ABC for parameter inference

The principle at the core of ABC is to approximate traditional Bayesian inference from a given dataset by simulations from the prior distribution. The simulated values are accepted or rejected according to the degree of proximity between the observed dataset y and a simulated one $y(\theta)$ thanks to a (usually normalized Euclidean) distance d. ABC relies on the operational assumptions that, while the likelihood is intractable, observations can be generated from the statistical model $f(\cdot \mid \theta)$ under consideration for a given value of the parameter θ .

The ABC resolution of the intractability issue with the likelihood is to produce a so-called *reference table*, recording a large number of datasets simulated from the prior predictive distribution, with density $f(y \mid \theta) \times \pi(\theta)$, and then extracting those that bring the simulations close enough to the actual sample. In most ABC implementations, for both computational and statistical efficiency reasons, the simulated $y^{(t)}$'s (t = 1, ..., N) are summarised through a dimension-reduction function $\eta : \mathscr{Y} \to \mathbb{R}^k$, often called a vector of k summary statistics. While the outcome of the ABC algorithm is then an approximation to the posterior distribution of θ given $\eta(y)$, rather than given the entire data y (Marin et al., 2012), arguments are to be found in the literature supporting the (ideal) choice of a summary statistic η of the same dimension as the parameter (Fearnhead and Prangle, 2012; Li and Fearnhead, 2015; Frazier et al., 2017). Algorithm 1 details how the *reference table* is constructed. The *reference table* will latter be used as a training dataset for the different RF methods explained below.

In practice, a *reference table* of size N is simulated, distances $(d(\eta(y), \eta(y^{(t)})))_{t=1,\dots,N}$ are computed and then given a tolerance proportion $0 < p_{\varepsilon} \le 1$, pairs $(\theta^{(t)}, \eta(y^{(t)}))$ within the p_{ε} range of lowest distances are selected. The parameter sample thus derived is deemed to approximate the posterior distribution $\pi(\theta \mid \eta(y))$.

Algorithm 1: Generation of a *reference table* from the prior predictive distribution $\pi(\theta) f(y \mid \theta)$

```
 \begin{array}{ll} \textbf{1 for } t \leftarrow 1 \textbf{ to } N \textbf{ do} \\ \textbf{2} & \text{Simulate } \boldsymbol{\theta}^{(t)} \sim \boldsymbol{\pi}(\boldsymbol{\theta}); \\ \textbf{3} & \text{Simulate } \boldsymbol{y}^{(t)} = (y_1^{(t)}, \dots, y_n^{(t)}) \sim f\left(\boldsymbol{y} \mid \boldsymbol{\theta}^{(t)}\right); \\ \textbf{4} & \text{Compute } \boldsymbol{\eta}(\boldsymbol{y}^{(t)}) = \{\boldsymbol{\eta}_1(\boldsymbol{y}^{(t)}), \dots, \boldsymbol{\eta}_k(\boldsymbol{y}^{(t)})\}; \\ \textbf{5 end} \end{array}
```

The method is asymptotically consistent in the sense that the true parameter behind the data can be exhibited when both the sample size and the number of simulations grow to infinity and the tolerance decreases to zero (Frazier et al., 2017). However, it suffers from two major drawbacks. First, to ensure a sufficient degree of reliability, the number N of simulations must be quite large. Hence, it may prove difficult to apply ABC on large or complex datasets since producing data may prove extremely costly. Second, the calibration of the ABC algorithm (i.e. a tolerance level indicating the separation of accepted from rejected simulated parameter values) is a critical step and impacts the resulting approximation (Marin et al., 2012; Blum et al., 2013). Since the justification of the method is doubly asymptotic, it is delicate if not impossible (Li and Fearnhead, 2015; Frazier et al., 2017) to optimally tune ABC for finite sample sizes. A third feature of major importance in this algorithm is that it requires selecting a vector of summary statistics that captures enough information from the observed and simulated data. For most problems, using the raw data to run the comparison is indeed impossible, if only because of the dimension of the data. Fearnhead and Prangle (2012) give a natural interpretation of the vector of summary statistics as an estimator of θ , but this puts a clear restriction on the dimension and nature of the components of $\eta(y)$.

Finally, it has to be noted that the original rejection ABC method, which can be interpreted as a K-nearest

neighbour method, has been recurrently improved by linear or non-linear regression strategies, mentioned in literature as adjusted local linear (Beaumont et al., 2002), ridge regression (Blum et al., 2013), and by methods based on adjusted neural networks (Blum and François, 2010). Instead of ridge regularization within local linear adjustment, Saulnier et al. (2017) propose to use a lasso regularization in order to select among the summary statistics. The obtained results are promising except when the summary statistics are highly correlated. In such cases, Saulnier et al. (2017) suggest to use random forests. We will only consider ridge regularization in the present work.

2.2 Random forest methodology

The random forest methodology (RF) of Breiman (2001) is pivotal in our proposal. We use Breiman's RF in a regression setting where a response variable $Y \in \mathbb{R}$ is explained by a vector of covariates $X = (X^{(1)}, \ldots, X^{(k)})$. A collection of N datasets, made of responses and associated covariates, is used to train a RF.

A given regression RF of size B is composed of B regression trees. A tree is a structure made of binary nodes, which are iteratively built from top to bottom until a stopping rule is satisfied. There are two types of nodes in such trees, the internal and terminal nodes, the latter being also called leaves. At an internal node, a binary rule of the form $X^{(j)} \leq s$ versus $X^{(j)} > s$ compares a covariate $X^{(j)}$ with a bound s. The result of the test divides the predictor space and the training dataset depending on this splitting rule into two parts in two new different nodes. When constructing the tree based on a training sample, the covariate index j and the splitting bound s are determined towards minimising a L^2 -loss criterion. The same covariate may be used multiple times for the choice of j at different levels of the tree construction. Splitting events stop when all the observations of the training dataset in a given node have the same covariates value, in that case the node becomes a leaf. Moreover, when a node has less than N_{\min} observations, the node also becomes a leaf, typically $N_{\min} = 5$ in the regression framework. Once the tree construction is complete, a value of the response variable is allocated to each tree leaf, corresponding to the average of the response values of the present datasets. For a given and an observed dataset that corresponds to a new covariate X, predicting the associated value of Y implies following the path of the binary rules. The outcome of the prediction is the allocated value of the leaf where this dataset ends after following this path.

The RF method consists in aggregating (or bagging) randomized regression trees. A large number of trees are trained on bootstrap samples of the training dataset and furthermore a subset of n_{try} covariates among the k available covariates are randomly considered at each split. The predicted value of a regression RF is determined by averaging the B predictions over its B tree components.

2.3 ABC parameter inference using random forest

2.3.1 Motivations and main principles

The particular choice of RF as a (non-parametric) estimation method in a regression setting is justified by the robustness of both random forests and quantile methods to "noise", that is, to the presence of irrelevant predictors, even when the proportion of such covariates amongst the entire set of proposed predictors is substantial (Marin et al., 2017). By comparison, the method of *K*-nearest neighbour classifiers lacks such characteristics (Biau et al., 2015). In the setting of building an ABC algorithm without preliminary selection of some summary statistics, our conjecture is that RF allows for the inclusion of an arbitrary and potentially large number of summary statistics in the derivation of the forest and therefore that it does not require the usual preliminary selection of summary statistics. When implementing this approach, we hence bypass

the selection of summary statistics and include a large collection of summary statistics, some or many of which being potentially poorly informative if not irrelevant for the regression. For earlier considerations on the selection of summary statistics, see Joyce and Marjoram (2008); Nunes and Balding (2010); Jung and Marjoram (2011); Fearnhead and Prangle (2012) and the review paper of Blum et al. (2013) where different dimension reduction techniques are compared.

A regression RF produces an expected predicted value for an arbitrary transform of θ , conditional on an observed dataset. This prediction is the output of a piece-wise constant function of the summary statistics. RF aggregates trees, partitions the feature space (here the space of summary statistics) in a way tuned to the forecast of a scalar output, i.e., a one dimensional functional of the parameter. This partition and prediction are done without requiring the definition of a particular distance on the feature space and is hence not dependant of any type of tolerance level. From an ABC perspective, each tree of a RF provides a partition of the summary space that is adapted for the forecasting of a scalar transform $h(\theta)$ of the parameter θ . In the following subsection we present how to compute quantities of interest in a context of parameter inference, thanks to the calculation of weights.

2.3.2 Calculation of weights and approximation of the posterior expectation

Assume we have now grown a RF made of B trees that predicts $\tau = h(\theta) \in \mathbb{R}$ using $\eta(y)$ and the training sample $(\eta(y^{(t)}), \tau^{(t)})_{t=1,\dots,N}$, where $\tau^{(t)} = h(\theta^{(t)})$. In the examples below, we will consider the case where h is the projection on a given coordinate of θ . Each of these B trees produces a partition of the space of summary statistics, with a constant prediction of the expected value of τ on each set of the partition. More precisely, given b-th tree in the forest, let us denote $n_b^{(t)}$ the number of times the pair $(\eta(y^{(t)}), \tau^{(t)})$ is repeated in the bootstrap sample that is used for building the b-th tree. Note that $n_b^{(t)}$ is equal to zero when the pair does not belong to the bootstrap sample. These pairs form the so-called out-of-bag sample of the b-th tree. Now, let $L_b(\eta(y))$ denote the leaf reached after following the path of binary choices given by the tree, which depends on the value of $\eta(y)$. The number of items of the bootstrap sample that fall in that leaf is

$$|L_b(\eta(y))| = \sum_{t=1}^N n_b^{(t)} \mathbf{1} \{ \eta(y^{(t)}) \in L_b(\eta(y)) \},$$

where 1 denotes the indicator function, and the mean value of τ of that leaf of the b-th tree is

$$\frac{1}{|L_b(\eta(y))|} \sum_{t=1}^{N} n_b^{(t)} \mathbf{1} \Big\{ \eta(y^{(t)}) \in L_b(\eta(y)) \Big\} \tau^{(t)}.$$

Averaging these B predictions of τ leads to an approximation of the posterior expected value of τ , also denoted mean value of τ , which can be written as follows:

$$\widetilde{\mathbb{E}}(\tau | \eta(y)) = \frac{1}{B} \sum_{t=1}^{N} \sum_{b=1}^{B} \frac{1}{|L_b(\eta(y))|} n_b^{(t)} \mathbf{1} \{ \eta(y^{(t)}) \in L_b(\eta(y)) \} \tau^{(t)}.$$

As exhibited by Meinshausen (2006), the above can be seen as a weighted average of τ along the whole training sample of size N made by the *reference table*. In fact, the weight of the t-th pair $(\eta(y^{(t)}), \tau^{(t)})$ given $\eta(y)$ is

$$w_t(\eta(y)) = \frac{1}{B} \sum_{b=1}^{B} \frac{1}{|L_b(\eta(y))|} n_b^{(t)} \mathbf{1} \Big\{ \eta(y^{(t)}) \in L_b(\eta(y)) \Big\}.$$

2.3.3 Approximation of the posterior quantile and variance

The weights $w_t(\eta(y))$ provide an approximation of the posterior cumulative distribution function (cdf) of τ given $\eta(y)$ as

$$\widetilde{F}(\tau \mid \boldsymbol{\eta}(y)) = \sum_{t=1}^{N} w_t(\boldsymbol{\eta}(y)) \mathbf{1} \{ \tau^{(t)} < \tau \}.$$

Posterior quantiles, and hence credible intervals, are then derived by inverting this empirical cdf, that is by plugging \tilde{F} in the regular quantile definition

$$\widetilde{\mathbb{Q}}_{\alpha}\{\tau \mid \eta(y)\} = \inf\left\{\tau : \widetilde{F}(\tau \mid \eta(y)) \ge \alpha\right\}.$$

This derivation of a quantile approximation is implemented in the R package quantregForest and the consistency of \tilde{F} is established in Meinshausen (2006).

An approximation of $Var(\tau \mid y)$ can be derived in a natural way from \widetilde{F} , leading to

$$\widehat{\operatorname{Var}}(\tau \mid \boldsymbol{\eta}(y)) = \sum_{t=1}^{N} w_t(\boldsymbol{\eta}(y)) \left(\tau^{(t)} - \sum_{u=1}^{N} w_u(\boldsymbol{\eta}(y)) \tau^{(u)} \right)^2.$$

2.3.4 Alternative variance approximation

Regarding the specific case of the posterior variance of τ , we propose a slightly more involved albeit manageable version of a variance estimate. Recall that, in any given tree b, some entries from the *reference table* are not included since each tree relies on a bootstrap sample of the training dataset. The out-of-bag simulations, i.e. unused in a bootstrap sample, can be exploited toward returning an approximation of $\mathbb{E}\{\tau \mid \eta(y^{(t)})\}$, denoted $\hat{\tau}_{\text{oob}}^{(t)}$. Indeed, given a vector of summary statistics $\eta(y^{(t)})$ of the training dataset, passing this vector down the ensemble of trees where it has not been used and mean averaging the associated predictions provide such an approximation. Since

$$\operatorname{Var}(\tau \mid \boldsymbol{\eta}(y)) = \mathbb{E}\left(\left[\tau - \mathbb{E}\left\{\tau \mid \boldsymbol{\eta}(y)\right\}\right]^2 \mid \boldsymbol{\eta}(y)\right),\,$$

we advocate applying the original RF weights $w_t(\eta(y))$ to the out-of-bag square residuals $(\tau^{(t)} - \hat{\tau}_{oob}^{(t)})^2$, which results in the alternative approximation

$$\widetilde{\operatorname{Var}}(\tau \mid \boldsymbol{\eta}(y)) = \sum_{t=1}^{N} w_t \{ \boldsymbol{\eta}(y) \} (\tau^{(t)} - \hat{\tau}_{\operatorname{oob}}^{(t)})^2.$$

Under the same hypotheses as Meinshausen (2006), this estimator converges when $N \to \infty$. Indeed, $\hat{\tau}_{\text{oob}}^{(t)}$ and $\sum_{t=1}^{N} w_t(\eta(y)) \tau^{(t)}$ tends to the same posterior expectation. Hence, the two variance estimators above mentioned are equivalent. A comparison between different variance estimators is detailed in the supplementary material. Owing to the results of this comparative study, we choose to use the above alternative variance estimator when presenting the results from two examples.

As a final remark, it is worth stressing that the approximation of the posterior covariance between a pair of parameters can be achieved thanks to a total of three RFs. The details of that statistical extension are presented in the Section 2 of the supplementary material.

2.3.5 A new R package for conducting parameter inferences using ABC-RF

When several parameters are jointly of interest, our recommended global strategy consists in constructing one independent RF for each parameter of interest and estimate from each RF several summary measurements of the posterior distribution (i.e. posterior expectation, quantiles and variance) of each parameter. Additional RFs might be constructed, however, if one is interested by estimating posterior covariance between pair of parameters.

Our R library aborf was initially developed for Bayesian model choice using ABC-RF as in Pudlo et al. (2016). The version 1.7 of aborf includes all the methods proposed in this paper to estimate posterior expectations, quantiles, variances (and covariances) of parameter(s). aborf version 1.7 is available on CRAN. We provide in the Section 3 of the supplementary material, a commented R code that will allow non expert users to run random forest inferences about parameters using the aborf package version 1.7.

3 Results from two examples

We illustrate the performances of our ABC-RF method for Bayesian parameter inference on a Normal toy-example and on a realistic population genetics example. In the first case and only in that case, approximations of posterior quantities can be compared with their exact counterpart. For both examples, we further compare the performances of our methodology with those of earlier ABC methods based on solely rejection, adjusted local linear (Beaumont et al., 2002), ridge regression (Blum et al., 2013), and adjusted neural networks (Blum and François, 2010).

For both illustrations, RFs were trained based on the functions of the R package ranger (Wright et al., 2017) with forests made of B = 500 trees, with $n_{\rm try} = k/3$ selected covariates (i.e. summary statistics) for split-point selection at each node, and with a minimum node size equals to 5 (Breiman, 2001, and see Section 3.3 Practical recommendations regarding the implementation of the ABC-RF algorithm). The other ABC methods in the comparison were based on the same *reference tables*, calling the corresponding functions in the R package abc (Csilléry et al., 2012; Nunes and Prangle, 2015) with its default parameters. ABC with neural network adjustment require the specification of the number of layers composing the neural network: we opted for the default number of layers in the R package abc, namely 10. A correction for heteroscedasticity is applied by default when considering regression adjustment approaches. Note that regression corrections are univariate for local linear and ridge regression as well as for RF, whereas neural network - by construction - performs multivariate corrections.

3.1 Normal toy example

We consider the hierarchical Normal mean model

$$y_i \mid \theta_1, \theta_2 \sim \mathcal{N}(\theta_1, \theta_2),$$

 $\theta_1 \mid \theta_2 \sim \mathcal{N}(0, \theta_2),$
 $\theta_2 \sim I\mathcal{G}(4, 3),$

where $I\mathscr{G}(\kappa,\lambda)$ denotes an inverse Gamma distribution with shape parameter κ and scale parameter λ . Let $y=(y_1,\ldots,y_n)$ be a *n*-sample from the above model. Given these conjugate distributions, the marginal

posterior distributions of the parameters θ_1 and θ_2 are closed-forms:

$$\theta_1 \mid y \sim \mathcal{T}\left(n+8, \frac{n\bar{y}}{n+1}, \frac{2\left(3+s^2/2 + n\bar{y}^2/(2n+2)\right)}{(n+1)(n+8)}\right)$$

$$\theta_2 \mid y \sim I\mathcal{G}\left(\frac{n}{2} + 4, 3 + \frac{s^2}{2} + \frac{n\bar{y}^2}{2n+2}\right),$$

where \bar{y} is the sample mean and $s^2 = \sum_{i=1}^n (y_i - \bar{y})^2$ the sum of squared deviations. $\mathcal{T}(v, a, b)$ denotes the general t distribution with v degrees of freedom (Marin and Robert, 2014).

From the above expressions and for a given sample y, it is straightforward to derive the exact values of $\mathbb{E}(\theta_1 \mid y)$, $\operatorname{Var}(\theta_1 \mid y)$, $\operatorname{E}(\theta_2 \mid y)$, $\operatorname{Var}(\theta_2 \mid y)$ and posterior quantiles for the two parameters. This provides us with a benchmark on which to assess the performances of ABC-RF. For the present simulation study, we opted for a *reference table* made of $N=10^4$ replicates of a sample of size n=10 and k=61 summary statistics. Those statistics included the sample mean, the sample variance, the sample median absolute deviation (MAD), all possible sums and products with these three elements resulting in eight new summary statistics and 50 additional independent (pure) noise variables that were generated from a uniform $\mathscr{U}_{[0,1]}$ distribution. The performances of our method were evaluated on an independent test table of size $N_{\text{pred}} = 100$, produced in the same way as the *reference table*. Current ABC methods (rejection, adjusted local linear, ridge and neural network) all depend on the choice of a tolerance level p_{ε} corresponding to the proportion of selected simulated parameters with lowest distances between simulated and observed summary statistics. On this example we consider a tolerance level of $p_{\varepsilon} = 0.01$ for ABC with rejection, and $p_{\varepsilon} = 0.1$ for the ABC methods with adjustment.

Figure 1 compares the exact values $\psi_1(y) = \mathbb{E}(\theta_1 \mid y)$, $\psi_2(y) = \mathbb{E}(\theta_2 \mid y)$, $\psi_3(y) = \text{Var}(\theta_1 \mid y)$ and $\psi_4(y) = \text{Var}(\theta_2 \mid y)$ with the estimates obtained from the ABC-RF approach. It shows that the proposed estimators have good overall performances for both $\psi_1(y)$ and $\psi_2(y)$, although one can see that $\psi_2(y)$ tends to be slightly overestimated. Our estimators perform less satisfactorily for both $\psi_3(y)$ and $\psi_4(y)$ but remain acceptable. Figure 2 shows furthermore that the quantile estimation are good for θ_1 if less accurate for θ_2 .

We then run an experiment to evaluate the precision of the marginal posterior approximation provided by ABC-RF for the parameter θ_1 , using two different test datasets and 40 independent *reference tables*. As exhibited in Figure S3, results are mixed. For one dataset, the fit is quite satisfactory, with the RF approximation showing only slightly fatter tails than the true posterior density distribution function (Figure S3; upper panel). For the other dataset, we obtain stronger divergence both in location and precision of the posterior density distribution function (Figure S3; lower panel).

Using the same *reference table*, we now compare our ABC-RF results with a set of four earlier ABC methods, namely, ABC methods based on straightforward rejection, adjusted local linear, ridge regression and adjusted neural networks. Normalized mean absolute errors are used to measure performances. The normalization being done by dividing the absolute error by the true value of the target. A normalized version offers the advantage of being hardly impacted when only a few observations get poorly predicted. Table 1 shows that the ABC-RF approach leads to results better than all other ABC methods for all quantities of interest. Expectations and quantiles are noticeably more accurately predicted. Figure 3 compares differences between estimated and true values of the posterior variances $\psi_3(y)$, $\psi_4(y)$. It shows the global underestimation associated with earlier ABC methods, when compared to ABC-RF, the latter only slightly overestimating the posterior variance. Finally, by looking at the width of the boxplots of Figure 3, we deduce that our ABC-RF estimations exhibits a lower estimation variability.

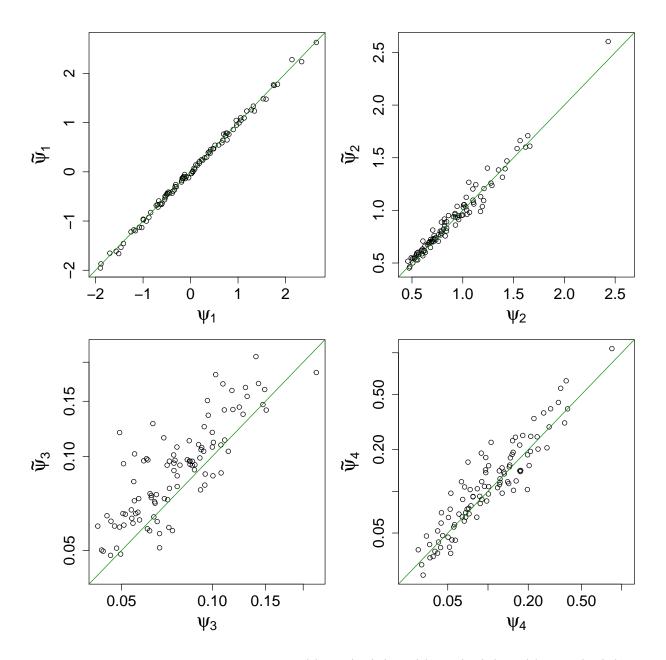


Figure 1: Scatterplot of the theoretical values $\psi_1(y) = \mathbb{E}(\theta_1 \mid y)$, $\psi_2(y) = \mathbb{E}(\theta_2 \mid y)$, $\psi_3(y) = \text{Var}(\theta_1 \mid y)$ and $\psi_4(y) = \text{Var}(\theta_2 \mid y)$ for the Normal model with their corresponding estimates $\tilde{\psi}_1$, $\tilde{\psi}_2$, $\tilde{\psi}_3$, $\tilde{\psi}_4$ obtained using ABC-RF. Variances are represented on a log scale.

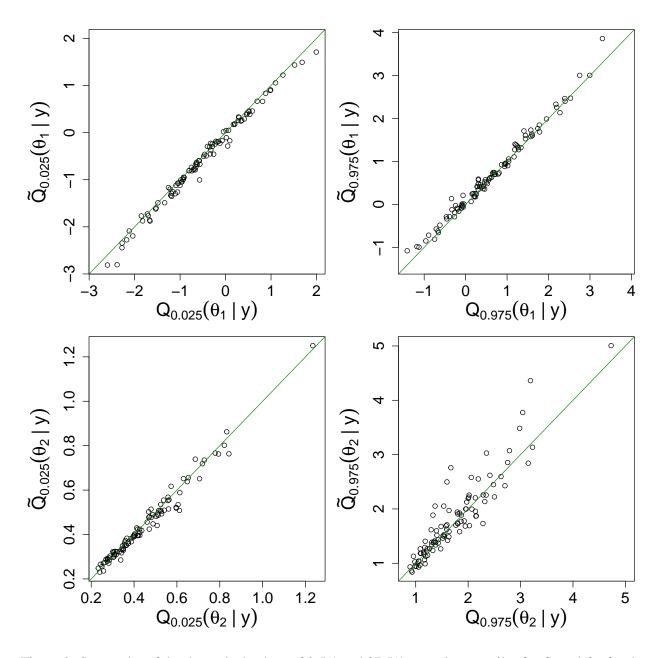


Figure 2: Scatterplot of the theoretical values of 2.5% and 97.5% posterior quantiles for θ_1 and θ_2 , for the Normal model with their corresponding estimates obtained using ABC-RF.

	RF	Reject	ALL	ARR	ANN
$\psi_1(y) = \mathbb{E}(\theta_1 \mid y)$	0.18	0.32	0.34	0.31	0.42
$\psi_2(y) = \mathbb{E}(\theta_2 \mid y)$	0.05	0.10	0.14	0.17	0.17
$\psi_3(y) = \operatorname{Var}(\theta_1 \mid y)$	0.25	2.21	0.70	0.69	0.48
$\psi_4(y) = \operatorname{Var}(\theta_2 \mid y)$	0.25	0.43	0.66	0.70	0.97
$Q_{0.025}(heta_1 y)$	0.34	1.61	0.69	0.84	0.50
$Q_{0.025}(heta_2 y)$	0.04	0.13	0.34	0.55	0.80
$Q_{0.975}(heta_1 y)$	0.25	1.35	0.53	0.70	0.60
$Q_{0.975}(\theta_2 y)$	0.10	0.14	0.20	0.20	0.42

Table 1: Comparison of normalized mean absolute errors (NMAE) of estimated quantities of interest obtained with ABC-RF and other ABC methodologies. RF, Reject, ALL, ARR and ANN stand for random forest (ABC-RF), rejection, adjusted local linear, adjusted ridge regression and adjusted neural network methods, respectively. The smallest NMAE values are in by bold characters.

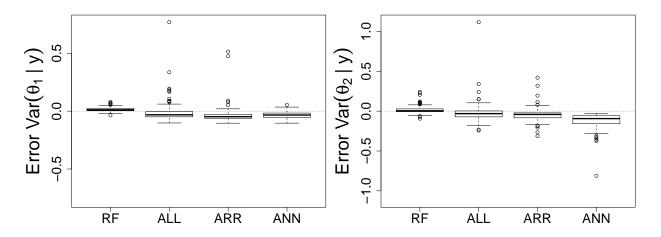


Figure 3: Boxplot comparison of the differences between our predictions for $Var(\theta_1 \mid y)$ and $Var(\theta_2 \mid y)$ and the corresponding true values, using ABC-RF and other ABC methods. RF, ALL, ARR and ANN notations as in the legend of Table 1). The closer to the y = 0 axis, the better the predictions. Boxplots above this axis imply overestimation of the predictions and below underestimation.

3.2 Human population genetics example

We now illustrate our methodological findings with the study of a population genetics dataset including 50000 single nucleotide polymorphic (SNP) markers genotyped in four human population samples (The 1000 Genomes Project Consortium, 2012; see details in Pudlo et al., 2016). The four populations include Yoruba (Africa; YRI), Han (East Asia; CHB), British (Europe; GBR) and American individuals of African ancestry (North America; ASW). The considered evolutionary model is represented in Figure 4. It includes a single out-of-Africa event with a secondarily split into one European and one East Asian population lineage and a recent genetic admixture of Afro-Americans with their African ancestors and with Europeans. The model was robustly chosen as most appropriate among a set of eight evolutionary models, when compared using ABC-RF for model choice in Pudlo et al. (2016).

We here focused our investigations on two parameters of interest in this model: (i) the admixture rate ra (i.e. the proportion of genes with a non-African origin) that describes the genetic admixture between individual of British and African ancestry in Afro-Americans individuals; and (ii) the ratio N2/Na between the ancestral effective population size Na and African N2 (in number of diploid individuals), roughly describing the increase of African population size in the past. Considering ratios of effective population sizes allows preventing identifiability issues of the model.

We used the software DIYABC v.2.0 (Cornuet et al., 2008, 2014) to generate a reference table of size 200000, with N = 199000 datasets being used as training dataset and $N_{pred} = 1000$ remaining as test datasets. RFs are built in the same way as for our Normal example and make use of the k = 112 summary statistics provided for SNP markers by DIYABC, (see Pudlo et al., 2016, and the Section 5 of the supplementary material).

Due to the complexity of this model, the exact calculation of any posterior quantity of interest is infeasible. To bypass this difficulty we compute NMAE using simulated parameters from the test table, rather than targeted posterior expectations. Here, 95% credible intervals (CI) are deduced from posterior quantile estimate of order 2.5% and 97.5%. Performances are measured via mean range and coverage, with coverage corresponding to the percentage of rightly bounded parameters. For example a 95% CI should provide coverage equal to 95% of the test table.

Figure 5 and Table 2 (see also Figure S4) illustrate the quality of the ABC-RF method when compared with ABC with either rejection, local linear, ridge or neural network adjustment (with logit transforms of the parameters for non rejection methods) using different tolerance levels (i.e., with tolerance proportion ranging from 0.005 to 1). We recall that considering the ABC rejection method with a tolerance equals to 1 is equivalent to work with the prior. Note that, due to memory allocation issues when using ABC method with adjusted ridge regression and a tolerance level of 1 on large *reference table*, we did not manage to recover results in this specific case.

Interesting methodological features can be observed in association with this example. ABC with rejection performs poorly in terms of NMAE and provides conservative and hence wide CIs (i.e., with coverage higher than the formal level). For ABC with adjustment, the lower the tolerance the lower the error (Table 2). The CI quality however highly suffers from low tolerance, with underestimated coverage (Figures 5 and S4). The smaller the tolerance value, the narrower the CI. Results for the ABC method with adjusted ridge regression seems however to be unstable for the parameter N2/Na depending on the considered level of tolerance. The ABC method using neural network and a tolerance level of 0.005 provides the lowest NMAE for both parameters of interest. The corresponding coverages are however underestimated, equal to 87.2% for ra and 81.6% for N2/Na, when 95% is expected (lower part of Figure 5 and Figure S4). Note that results with this method can be very time consuming to obtain when the tolerance level and the number of layers are large. The ABC-RF method provides an appealing trade-off between parameter estimation

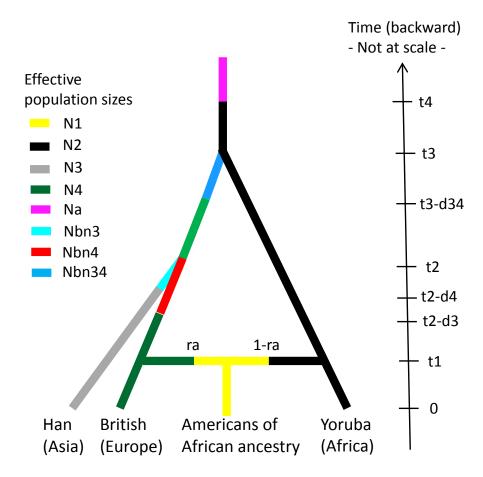


Figure 4: Evolutionary model of four human populations considered for Bayesian parameter inference using ABC-RF. The prior distributions of the demographic and historical parameters used to simulate SNP datasets are as followed: Uniform[100; 10 000] for the split times t2 and t3 (in number of generations), Uniform[1; 30] for the admixture time t1, Uniform[0.05; 0.95] for the admixture rate ra (proportion of genes with a non-African origin), Uniform[1000; 100 000] for the stable effective population sizes N1, N2, N3, N4 and N34 (in number of diploid individuals), Uniform[5; 500] for the bottleneck effective population sizes Nbn3, Nbn4, and Nbn34, Uniform[5; 500] for the bottleneck durations d3, d4, and d34, Uniform[100; 10 000] for both the ancestral effective population size Na and t4 the time of change to Na. Conditions on time events were t4>t3>t2. See Pudlo et al. (2016) for details. Regarding the genetic model, we simulated biallelic polymorphic SNP datasets using the algorithm proposed by Hudson (2002) (cf "-s 1" option in the program ms associated to Hudson (2002)). This coalescent-based algorithm provides the simulation efficiency and speed necessary in the context of ABC, where large numbers of simulated datasets including numerous (statistically independent) SNP loci have to be generated (see Supplementary Appendix S1 of Cornuet et al. (2014) for additional comments on Hudson?s algorithm).

Method	Tolerance level	ra NMAE	N2/Na NMAE
RF	NA	0.018	0.053
RF^*	NA	0.019	0.053
Reject	0.005	0.151	0.355
Reject	0.01	0.178	0.454
Reject	0.1	0.322	1.223
Reject	0.4	0.574	2.025
Reject	1	0.856	4.108
ALL	0.005	0.028	0.166
ALL	0.01	0.028	0.249
ALL	0.1	0.035	0.139
ALL	0.4	0.044	0.170
ALL	1	0.062	0.209
ARR	0.005	0.027	0.220
ARR	0.01	0.027	0.317
ARR	0.1	0.035	0.140
ARR	0.4	0.044	0.163
ARR	1	_	_
ANN	0.005	0.007	0.037
ANN	0.01	0.007	0.038
ANN	0.1	0.013	0.064
ANN	0.4	0.016	0.123
ANN	1	0.025	0.095

Table 2: Comparison of normalized mean absolute errors (NMAE) for the estimation of the parameters ra and N2/Na using ABC-RF (RF) and ABC with rejection (Reject), adjusted local linear (ALL) or ridge regression (ARR) or neural network (ANN) with various tolerance levels for Reject, ALL, ARR and ANN. NA stands for not appropriate. The smallest NMAE values are in by bold characters. NA stands for not appropriate. RF* refers to results obtained using ABC-RF when adding 20 additional independent noise variables generated from a uniform $\mathcal{U}_{[0,1]}$ distribution. RF refers to results without noise variables.

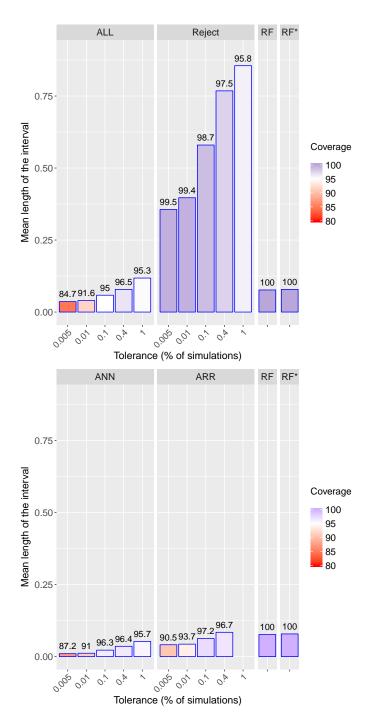


Figure 5: Range and coverage comparison of approximate 95% credible intervals on the admixture parameter ra (Figure 4) obtained with ABC-RF (RF) and with earlier ABC methods: rejection (Reject), adjusted local linear (ALL) or ridge regression (ARR) or neural network (ANN) with various tolerance levels for Reject, ALL, ARR and ANN. Coverages values are specified by bar colors and superimposed values. Heights indicate CI mean lengths. Results for ALL, Reject and RF are presented in the upper figure whereas those for ANN, ARR and RF are in the lower figure. See Figure S4 for a similar representation of results for the parameter N2/Na. RF* refers to results obtained using ABC-RF when adding 20 additional independent noise variables generated from a uniform $\mathcal{U}_{[0,1]}$ distribution. RF refers to results without noise variables.

quality (ABC-RF is the method with the second lowest NMAE values in Table 2) and slightly conservative CIs (Figures 5 and S4). Similar results and methodological features were observed when focusing on the 90% CI (results not shown). It is also worth stressing that not any calibration of any kind of a tolerance level parameter are needed with ABC-RF, which is an important plus for this method. On the opposite, earlier ABC methods require calibration to optimize their use, such calibration being time consuming when different levels of tolerance are used.

For the observed dataset used in this study, posterior expectations and quantiles of the parameters of interest ra and N2/Na are reported in Table 3. Expectation and CI values substantially vary for both parameters, depending on the method used. The impact of the tolerance levels is noteworthy for both the rejection and local linear adjustment ABC methods. The posterior expectation of ra obtained using ABC-RF was equal to 0.221 with a relatively narrow associated 95% CI of [0.112,0.287]. The latter estimation lays well within previous estimates of the mean proportion of genes of European ancestry within African American individuals, which typically ranged from 0.070 to 0.270 — with most estimates around 0.200 —, depending on individual exclusions, the population samples and sets of genetic markers considered, as well as the evolutionary models assumed and inferential methods used (reviewed in Bryc et al., 2015). Interestingly, a recent genomic analysis using a conditional random field parametrized by random forests trained on reference panels (Maples et al., 2013) and 500000 SNPs provided a similar expectation value of ra for the same African American population ASW (i.e. ra = 0.213), with a somewhat smaller 95% CI (i.e. [0.195,0.232]), probably due to the ten times larger number of SNPs in their dataset (Baharian et al., 2016).

The posterior expectation of N2/Na obtained using ABC-RF was equal to 4.508 with a narrow associated 95% CI of [3.831,5.424]. Such values point to the occurrence of the substantial ancestral demographic and geographic expansion that is widely illustrated in previous Human population genetics studies, including African populations (e.g. Henn et al., 2012). Although our modeling setting assumes a naïve abrupt change in effective population sizes in the ancestral African population, the equivalent of N2/Na values inferred from different methods and modeling settings fit rather well with our own posterior expectations and quantiles for this parameter (e.g. Schiffels and Durbin, 2014).

In contrast to earlier ABC methods, the RF approach is deemed to be mostly insensitive to the presence of covariates whose the distributions does not depend on the parameter values (i.e. ancillary covariates) (e.g. Breiman, 2001; Marin et al., 2012). To illustrate this feature, we have added 20 additional independent noise variables generated from a uniform $\mathcal{U}_{[0,1]}$ distribution (results designated by RF*) in the *reference table* generated for the present Human population genetics example. We found that the presence of such noise covariates do not impact the results in terms of NMAE, coverage and only slightly on parameter estimation for the observed dataset (Tables 2 and 3, and Figures 5 and S4). For the rest of the article, no noise variables were used.

3.3 Practical recommendations regarding the implementation of the ABC-RF algorithm

We mainly consider in this section two important practical issues, namely the choice of the number of simulations (N) in the *reference table* and of the number of trees (B) in the random forest. For sake of simplicity and concision, we focus our recommendations on the above human population genetics example (subsection 3.2). We stress here that, although not generic, our recommendations fit well with other examples of complex model settings that we have analysed so far (results not shown). We also stress that for simpler model settings substantially smaller N and B values were sufficient to obtain good results. Finally, we provide practical comments about the main sources of variabilities in inferences typical of the ABC-RF methodology.

						_
Method	Tol. level	Expectation	$Q_{0.025}$	$Q_{0.05}$	$Q_{0.95}$	$Q_{0.975}$
RF	NA	0.221	0.112	0.134	0.279	0.287
RF^*	NA	0.225	0.112	0.142	0.282	0.290
Reject	0.005	0.223	0.061	0.069	0.364	0.389
Reject	0.01	0.220	0.060	0.070	0.389	0.418
Reject	0.1	0.276	0.062	0.074	0.511	0.543
Reject	0.4	0.388	0.068	0.086	0.739	0.791
Reject	1	0.502	0.073	0.095	0.906	0.928
ALL	0.005	0.278	0.219	0.229	0.322	0.337
ALL	0.01	0.257	0.232	0.238	0.274	0.278
ALL	0.1	0.207	0.170	0.171	0.233	0.237
ALL	0.4	0.194	0.144	0.152	0.233	0.241
ALL	1	0.196	0.115	0.126	0.278	0.299
ARR	0.005	0.260	0.252	0.254	0.265	0.266
ARR	0.01	0.252	0.239	0.242	0.260	0.262
ARR	0.1	0.211	0.171	0.178	0.239	0.244
ARR	0.4	0.196	0.140	0.149	0.241	0.251
ARR	1	_	_	_	_	_
ANN	0.005	0.227	0.221	0.223	0.232	0.234
ANN	0.01	0.226	0.219	0.221	0.231	0.233
ANN	0.1	0.228	0.217	0.220	0.236	0.239
ANN	0.4	0.232	0.216	0.221	0.242	0.248
ANN	1	0.206	0.183	0.187	0.227	0.233

N	7	ľ	N	

Method	Tol. level	Expectation	$Q_{0.025}$	$Q_{0.05}$	$Q_{0.95}$	$Q_{0.975}$
RF	NA	4.508	3.831	3.959	5.153	5.424
RF^*	NA	4.594	3.821	3.910	5.241	6.552
Reject	0.005	6.282	2.937	3.223	10.086	11.337
Reject	0.01	6.542	2.746	3.116	10.837	11.852
Reject	0.1	8.001	2.131	2.574	15.690	18.531
Reject	0.4	11.605	1.795	2.331	28.011	38.532
Reject	1	23.483	0.672	1.185	84.649	147.657
ALL	0.005	30.041	1.256	1.879	83.369	174.340
ALL	0.01	9.289	3.946	4.586	16.686	20.361
ALL	0.1	8.235	5.736	5.995	11.573	12.719
ALL	0.4	10.752	4.588	4.996	21.656	27.300
ALL	1	7.222	5.684	5.829	9.631	10.475
ARR	0.005	10.528	4.395	5.677	19.224	22.722
ARR	0.01	8.264	5.020	5.485	12.544	13.313
ARR	0.1	8.394	5.643	5.948	12.075	13.313
ARR	0.4	10.802	6.113	6.505	17.487	20.511
ARR	1	_	_	_	_	_
ANN	0.005	5.746	5.512	5.563	5.937	5.982
ANN	0.01	6.148	5.883	5.934	6.353	6.420
ANN	0.1	25.921	23.857	24.250	27.672	28.133
ANN	0.4	8.515	7.652	7.810	9.147	9.436
ANN	1	7.021	5.692	5.856	8.677	9.370

Table 3: Estimation of the two parameters of interest ra and N2/Na for the observed human population genetics dataset using ABC-RF (RF), and ABC with rejection (Reject), adjusted local linear (ALL) or ridge regression (ARR) or neural network (ANN) with various tolerance levels (Tol. level) for Reject, ALL, ARR and ANN. NA stands for not appropriate. RF* refers to results obtained using ABC-RF when adding 20 additional independent noise variables generated from a uniform $\mathcal{U}_{[0,1]}$ distribution. RF refers to results without noise variables.

NMAE							
$N (\times 10^3)$	10	25	50	75	100	150	199
ra	0.028	0.023	0.021	0.020	0.019	0.018	0.018
N2/Na	0.080	0.067	0.059	0.057	0.055	0.053	0.053
OOB MSE							
$N (\times 10^3)$	10	25	50	75	100	150	199
$ra (\times 10^{-4})$	1.670	1.176	0.914	0.823	0.745	0.695	0.664
N2/Na $(\times 10^3)$	0.194	0.179	0.143	0.125	0.115	0.111	0.110

Table 4: Comparison of normalized mean absolute errors (NMAE) and out-of-bag mean squared errors (OOB MSE) for the estimation of the parameters ra and N2/Na obtained with ABC-RF, using different *reference table* sizes (N). We use the test table mentioned in subsection 3.2. The number of trees in the RF is 500.

$N (\times 10^3)$	10	25	50	75	100	150	199
ra expectation	0.231	0.222	0.224	0.223	0.222	0.223	0.221
ra $Q_{0.025}$	0.097	0.095	0.102	0.104	0.106	0.109	0.112
ra $Q_{0.975}$	0.317	0.309	0.305	0.305	0.289	0.292	0.287
N2/Na expectation	4.538	4.588	4.652	4.530	4.475	4.483	4.508
N2/Na Q _{0.025}	3.651	3.679	3.782	3.802	3.751	3.840	3.831
N2/Na Q _{0.975}	6.621	6.221	6.621	5.611	5.555	5.315	5.424

Table 5: Estimation of the parameters ra and N2/Na for the observed population genetics dataset with ABC-RF, using different *reference table* sizes (*N*). The number of trees in the RF is 500.

Reference table size — We consider a reference table made of $N = 199\,000$ simulated datasets. However, Table 4 shows a negligible decrease of NMAE when using $N = 100\,000$ to $N = 199\,000$ datasets. Table 5 also exhibits small variations between predictions on the observed dataset, especially for $N \ge 7500$. The level of variation thus seems to be compatible with the random variability of the RF themselves. Altogether, using a reference table including $100\,000$ datasets seems to be a reasonable default choice. It is worth stressing that the out-of-bag mean squared error can be easily retrieved and provides a good indicator of the quality of the RF without requiring the simulation of a (small size) secondary test table, which can hence be estimated at a low computational cost (Table 4).

Number of trees — A forest including 500 trees is a default choice when building RFs, as this provides a good trade-off between computation efficiency and statistical precision (Breiman, 2001; Pudlo et al., 2016). To evaluate whether or not this number is sufficient, we recommend to compute the out-of-bag mean squared error depending on the number of trees in the forest for a given *reference table*. If 500 trees is a satisfactory calibration, one should observe a stabilization of the error around this value. Figure 6 illustrates this representation on the human population genetics example and points to a negligible decrease of the error after 500 trees. This graphical representation is produced via our R package abcrf.

Minimum node size (maximum leaf size) — We recall that splitting events during a tree construction stop when a node has less than N_{\min} observations, in that case, the node becomes a leaf. Note that the higher N_{\min}

NMAE

$N_{ m min}$	1	2	3	4	5	10	20	50	100	200	500
ra	0.019	0.019	0.019	0.019	0.019	0.019	0.020	0.021	0.023	0.027	0.033
N2/Na	0.054	0.055	0.054	0.055	0.055	0.055	0.055	0.058	0.062	0.068	0.082
				C	OB MS	Е					
N_{\min}	1	2	3	4	5	10	20	50	100	200	500
$ra (\times 10^{-4})$	0.745	0.739	0.744	0.739	0.745	0.760	0.783	0.925	1.129	1.480	2.280
$N2/N_2 (\times 10^3)$	0.114	0.116	0.115	0.115	0.115	0.116	0.110	0.131	0.153	0.183	0.252

Table 6: Comparison of normalized mean absolute errors (NMAE) and out-of-bag mean squared errors (OOB MSE) for the estimation of the parameters ra and N2/Na obtained with ABC-RF, using different minimum node sizes (N_{\min}). We use the reference table of size $N = 100\,000$ and the test table mentioned in subsection 3.2. The number of trees in the RF is 500.

the quicker RF treatments. In all RF treatments presented here, we used the default size $N_{\min} = 5$. Table 6 illustrates the influence of N_{\min} on the human population genetics example and highlights a negligible decrease of the error for N_{\min} lower than 5.

Finally, we see no reason to change the number of summary statistics sampled at each split n_{try} within a tree, which is traditionally chosen as k/3 for regression when k is the total number of predictors (Breiman, 2001).

Variability in the ABC-RF methodology — The ABC-RF methodology is associated with different sources of variabilities the user should be aware of. Using a simulated *reference table* is the main source, RF being the second. Indeed, predicting quantities of interest for the same test dataset with two different *reference tables* of equal size N will result in slightly different estimates. This variation has been previously highlighted in Figure S3 dealing with the analysis of the Normal toy example. We recall RF are composed of trees trained on bootstrap samples, each one considering $n_{\rm try}$ covariates randomly selected amongst the k available at each split. This random aspects of RF results in variability. In practice, a good user habit should be to run ABC-RF more than once on different training datasets to ensure that the previously mentioned variabilities are negligible. If this variability is significant, we recommend considering a *reference table* of higher size.

4 Discussion

This paper introduces a novel approach to parameter estimation in likelihood-free problems, relying on the machine-learning tool of regression RF to automatise the inclusion of summary statistics in ABC algorithms. Our simulation experiments demonstrate several advantages of our methodological proposal compared with earlier ABC methods.

While using the same *reference table* and test dataset for all compared methods, our RF approach appears to be more accurate than previous ABC solutions. Approximations of expectations are quite accurate, while posterior variances are only slightly overestimated, which is an improvement compared with other approaches that typically underestimate these posterior variances. The performances for covariance approximation are quite encouraging as well, although the method is still incomplete and need further developments

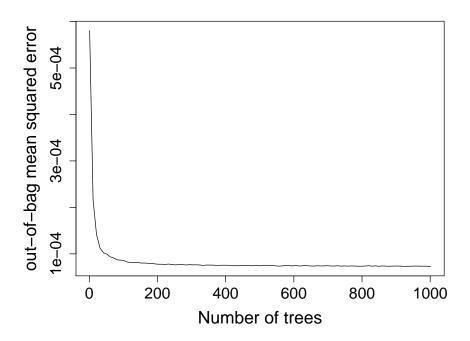


Figure 6: Relations between the number of trees in the forest and the ABC-RF out-of-bag mean squared errors, for a reference table of size N = 100000 in the human population genetics example.

on this particular point (more details are given in Section 2 of the supplementary material). We found that quantile estimations depend on the corresponding probability and we believe this must be related to the approximation error of the posterior cumulative function $F(x \mid \eta(y))$. More specifically, we observed that upper quantiles may be overestimated, whereas lower quantiles may be underestimated (Figure 2), indicating fatter tails in the approximation. Hence, credible intervals produced by the RF solution may be larger than the exact ones. However from a risk assessment point of view, this overestimation aspect clearly presents less drawbacks than underestimation of credible intervals. Altogether, owing to the various models and datasets we analysed, we argue that ABC-RF provides a good trade-off in terms of quality between parameter estimation of point estimators (e.g. expectation, median or variance) and credible interval coverage. A comparison of computing times is given in Section 8 of the supplementary material.

Throughout our experiments, we found that, contrary to earlier ABC methods, the RF approach is mostly insensitive to the presence of covariates whose the distributions does not depend on the parameter values (ancillary covariates). Therefore, we argue that the RF method can deal with a very large number of summary statistics, bypassing any form of pre-selection of those summaries. Interestingly, the property of ABC-RF to extract and adaptively weight information carried by each of the numerous summary statistics proposed as explanatory variables can be represented by graphs, showing the relative contribution of summary statistics in ABC-RF estimation for each studied parameter (see Section 7 of the supplementary material for details).

As an alternative, Papamakarios and Murray (2016) propose to approximate the whole posterior distribution by using Mixture Density Networks (MDN, Bishop, 1994). The MDN strategy consists in using Gaussian mixture models with parameters calibrated thanks to neural networks. The strategy of Papamakarios and Murray (2016) is to iteratively learn an efficient proposal prior (approximating the posterior distribution), then to use this proposal to train the posterior, both steps making use of MDN. This strategy can be easily applied when the prior is uniform or Gaussian, but other prior choices can involve difficulties. This is because in such cases, it might be difficult to simulate from the corresponding proposal. The approx-

imation accuracy of the posterior as a Gaussian mixture model depends of the number of components and the number of hidden layers of the networks. Those two parameters require calibration. Finally, by using MDN, one loses the contribution of summary statistics provided by RF and thus some useful interpretation elements. Despite these remarks, this promising method remains of interest and is worth mentioning.

The RF method focuses on unidimensional parameter inference. Multi-objective random forest (Kocev et al., 2007) could be a solution to deal with multidimensional parameter using RF. However, our attempts based on the later methodology were so far unfruitful (results not shown). An alternative approach could be based on using the RF strategies to approximate some conditionals distributions and then recover the joined posterior using either a Gibbs sampler (based on approximated full conditionals) or a Russian rule decompositions to which a product of embedded full conditionals is associated. We are presently comparing the two strategies on simulated datasets.

In population genetics, which historically corresponds to the field of introduction of ABC methods, next generation sequencing technologies result in large genome-wide datasets that can be quite informative about the demographic history of the genotyped populations. Several recently developed inferential methods relying on the observed site frequency spectrum appear particularly well suited to accurately characterizing the complex evolutionary history of invasive populations (Gutenkunst et al., 2009; Excoffier et al., 2013). Because of the reduced computational resources demanded by ABC-RF and the above-mentioned properties of the method, we believe that ABC-RF can efficiently contribute to the analysis of massive SNP datasets, including both model choice (Pudlo et al., 2016) and Bayesian inference about parameters of interest. More generally, the method should appeal to all scientific fields in which large datasets and complex models are analysed using simulation-based methods such as ABC (e.g. Beaumont, 2010; Sisson et al., 2017).

Acknowledgements

The use of random forests was suggested to J.-M.M. and C.P.R. by Bin Yu during a visit at CREST, Paris. We are very grateful to her for this conversation and also to Nicolai Meinshausen for some very helpful exchanges on the quantile regression forest method. We thank Dennis Prangle and Michael Blum for their very helpful suggestions, helping us to improve the quality of that paper.

Funding

A.E. acknowledges financial support by the National Research Fund ANR (France) through the European Union program ERA-Net BiodivERsA (project EXOTIC) and the project ANR-16-CE02-0015-01 (SWING), and the INRA scientific department SPE (AAP-SPE 2016). C.P.R. is supported by an IUF 2016–2021 senior grant.

References

- Arenas, M. et al. (2015) CodABC: A Computational Framework to Coestimate Recombination, Substitution, and Molecular Adaptation Rates by Approximate Bayesian Computation. *Molecular Biology and Evolution*, 32(4):1109–1112.
- Baharian, S. et al. (2016) The Great Migration and African-American Genomic Diversity. *PLOS Genetics*, 12(5):e1006059.
- Beaumont, M. A. et al. (2002) Approximate Bayesian Computation in Population Genetics. *Genetics*, 162(4):2025–2035.
- Beaumont, M. A. (2008) Joint determination of topology, divergence time and immigration in population trees. In Simulations, Genetics and Human Prehistory, pages 134–154. Eds. Matsumura S., Forster P. and Renfrew C. *Cambridge: (McDonald Institute Monographs), McDonald Institute for Archaeological Research*, 2008.
- Beaumont, M. A. (2010) Approximate Bayesian Computation in Evolution and Ecology. *Annual Review of Ecology, Evolution, and Systematics*, 41:379–406.
- Biau, G. et al. (2015) New Insights into Approximate Bayesian Computation. *Annales de l'Institut Henri Poincaré B, Probability and Statistics*, 51(1):376–403.
- Bishop, C. M. (1994) Mixture Density Networks. Tech. Rep. NCRG/94/004, Neural Computing Research Group, Aston University.
- Blum, M. G. B. and François, O. (2010) Non-linear regression models for Approximate Bayesian Computation. *Statistics and Computing*, 20(1):63–73.
- Blum, M. G. B. et al. (2013) A comparative review of dimension reduction methods in Approximate Bayesian Computation. *Statistical Science*, 28(2):189–208.
- Breiman, L. (2001) Random forests. *Machine Learning*, 45(1):5–32.
- Bryc, K. et al. (2015) The Genetic Ancestry of African Americans, Latinos, and European Americans across the United States. *The American Journal of Human Genetics*, 96(1):37-53.
- Chan, Y. L. et al. (2014) Detecting Concerted Demographic Response across Community Assemblages Using Hierarchical Approximate Bayesian Computation. *Molecular Biology and Evolution*, 31(9):2501–2515.
- Chernozhukov, V. and Hong, H. (2003) An MCMC approach to classical estimation. *Journal of Econometrics*, 115(2):293–346.
- Csilléry, K. et al. (2010) Approximate Bayesian Computation (ABC) in practice. *Trends in Ecology and Evolution*, 25(7):410–418.
- Csilléry, K. et al. (2012) abc: an R package for approximate Bayesian computation (ABC). *Methods in Ecology and Evolution*, 3(3):475–479.

- Cornuet, J.-M. et al. (2008) Inferring population history with DIY ABC: a user-friendly approach to approximate Bayesian computation. *Bioinformatics*, 24(23), 2713-2719.
- Cornuet, J.-M. et al. (2014) DIYABC v2.0: a software to make approximate Bayesian computation inferences about population history using single nucleotide polymorphism, DNA sequence and microsatellite data. *Bioinformatics*, 30:(8), 1187-1189.
- Excoffier, L. et al. (2013) Robust Demographic Inference from Genomic and SNP Data. *PLOS Genetics*, 9(10):e1003905.
- Fearnhead, P. and Prangle, D. (2012) Constructing summary statistics for approximate Bayesian computation: semi-automatic approximate Bayesian computation. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 74(3):419–474.
- Frazier, D. T. et al. (2017) Asymptotic Properties of Approximate Bayesian Computation. *arXiv* 1609.06903.
- Gourieroux, C. et al. (1993) Indirect inference. Journal of Applied Econometrics, 8:85-118.
- Gutenkunst, R. N. et al. (2009) Inferring the Joint Demographic History of Multiple Populations from Multidimensional SNP Frequency Data. *PLOS Genetics*, 5(10):e1000695.
- Henn, B. M. et al. (2012) The great human expansion. Proc Notl Acad Sci USA, 109(44):17758-17764.
- Hudson, R. (2002) Generating samples under a Wright-Fisher neutral model of genetic variation. *Bioinformatics*, 18(2):337-338.
- Joyce, P. and Marjoram, P. (2008) Approximately sufficient statistics and Bayesian computation. *Statistical Application in Genetics and Molecular Biology*, 7(1):Article 26.
- Jung, H. and Marjoram, P. (2011) Choice of Summary Statistics Weights in Approximate Bayesian Computation. *Statistical Application in Genetics and Molecular Biology*, 10(1):Article 45.
- Kocev, D. et al. (2007) Ensembles of Multi-Objective Decision Trees. In Machine Learning: ECML 2007. Lecture Notes in Computer Science, vol 4701. pages 624–631. Eds. Kok J. N., Koronacki J., Mantaras R. L., Matwin S., Mladenič D. and Skowron A. *Springer, Berlin, Heidelberg*, 2007.
- Li, W. and Fearnhead, P. (2015) On the Asymptotic Efficiency of ABC Estimators. arXiv, 1506.03481.
- Lindsay, B. (1988) Composite likelihood methods. Contemporary Mathematics, 80:221–239.
- Maples, B. K. et al. (2013). RFMix: A discriminative modeling approach for rapid and robust local-ancestry inference. *The American Journal of Human Genetics*, 93(2):278-288.
- Marin, J.-M. et al. (2012) Approximate Bayesian computational methods. *Statistics and Computing*, 22(6): 1167–1180.
- Marin, J.-M. and Robert, C. P. (2014) Bayesian Essentials with R. Springer.
- Marin, J.-M. et al. (2017) Likelihood-free Model Choice. In Handbook of Approximate Bayesian Computation. Eds. Sisson S. A., Fan Y. and Beaumont M. A. *Chapman and Hall/CRC Press*, in press.

- Meinshausen, N. (2006) Quantile Regression Forests. Journal of Machine Learning Research, 7:983–999.
- Nunes, M. A. and Balding, D. J. (2010) On Optimal Selection of Summary Statistics for Approximate Bayesian Computation. *Statistical Application in Genetics and Molecular Biology*, 9(1):Article 34.
- Nunes, M. and Prangle, D. (2015) abctools: An R Package for Tuning Approximate Bayesian Computation Analyses. *The R Journal*, 7(2):189–205.
- Papamakarios, G. and Murray, I. (2016) Fast ε -free Inference of Simulation Models with Bayesian Conditional Density Estimation. In Advances in Neural Information Processing Systems 29, pages 1028–1036. Eds. Lee D. D., Sugiyama M., Luxburg U. V., Guyon I. and Garnett R. *Curran Associates, Inc.*, 2016.
- Pritchard, J. K. et al. (1999) Population growth of human Y chromosomes: a study of Y chromosome microsatellites. *Molecular Biology and Evolution*, 16(12):1791–1798.
- Pudlo, P. et al. (2016) Reliable ABC model choice via random forests. Bioinformatics, 32(6):859-866.
- Saulnier, E. et al. (2017) Inferring epidemiological parameters from phylogenies using regression-ABC: A comparative study. *PLOS Computational Biology*, 13(3): e1005416.
- Schiffels, S. and Durbin, R. (2014) Inferring human population size and separation history from multiple genome sequences. *Nature Genetics*, 46:919-925.
- Sisson, S. A. et al. (2017) Handbook of Approximate Bayesian Computation. *Chapman & Hall/CRC Press*, in press.
- Tavaré, S. et al. (1997) Inferring Coalescence Times from DNA Sequence Data. *Genetics*, 145(2):505–518.
- The 1000 Genomes Project Consortium. (2012) An integrated map of genetic variation from 1,092 human genomes. *Nautre*, 491:56-65.
- Theunert, C. et al. (2012) Inferring the History of Population Size Change from Genome-Wide SNP Data. *Molecular Biology and Evolution*, 29(12):3653–3667.
- Toni, T. et al. (2009) Approximate Bayesian computation scheme for parameter inference and model selection in dynamical systems. *Journal of the Royal Society Interface*, 6(31):187–202.
- Varin, C. et al. (2011) An overview of composite likelihood methods. *Statistica Sinica*, 21(1):5–42.
- Wright, M. N. and Ziegler, A. (2017) ranger: A Fast Implementation of Random Forests for High Dimensional Data in C++ and R. *Journal of Statistical Software*, 77(1):, 1-17.

Supplementary Information ABC random forests for Bayesian parameter inference

1 Comparing three methods of variance estimation of parameters

We here compare three methods to estimate posterior variance of a parameter transformation of interest τ using ABC-RF. Two of them have already been explained in the main text (i.e. method 1 and 3 below).

• **Method 1**: One reuses the original random forest (RF) weights $w_t(\eta(y))$ to the out-of-bag square residuals $(\tau^{(t)} - \hat{\tau}_{\text{oob}}^{(t)})^2$, giving the variance estimator

$$\widetilde{\operatorname{Var}}(\tau \mid \boldsymbol{\eta}(y)) = \sum_{t=1}^{N} w_{t}(\boldsymbol{\eta}(y)) (\tau^{(t)} - \hat{\tau}_{\operatorname{oob}}^{(t)})^{2}.$$

• **Method 2**: A similar estimator can be obtained by building a new RF thanks to the training sample $(\eta(y^{(t)}), (\tau^{(t)} - \tau_{\text{oob}})^2)_{t=1,\dots,N}$, resulting in the estimator

$$\operatorname{Var}^{\#}(\tau \mid \boldsymbol{\eta}(y)) = \sum_{t=1}^{N} \widetilde{w_{t}}(\boldsymbol{\eta}(y)) (\tau^{(t)} - \hat{\tau}_{\operatorname{oob}}^{(t)})^{2},$$

where $\widetilde{w}_t(\eta(y))$ is the computed weights of this newly trained RF. This estimator is based on the expression of the posterior variance as a conditional expectation:

$$\operatorname{Var}(\tau \mid \boldsymbol{\eta}(y)) = \mathbb{E}\left(\left[\tau - \mathbb{E}(\tau \mid \boldsymbol{\eta}(y))\right]^{2} \mid \boldsymbol{\eta}(y)\right)$$

and the fact that such a RF is able to estimate this posterior expectation. This approach is more expensive due to the additional RF requirement.

• **Method 3**: The variance estimator is based on the cumulative distribution function (cdf) approximation,

$$\widehat{\operatorname{Var}}(\tau \mid \boldsymbol{\eta}(y)) = \sum_{t=1}^{N} w_t(\boldsymbol{\eta}(y)) \left(\tau^{(t)} - \sum_{u=1}^{N} w_u(\boldsymbol{\eta}(y)) \tau^{(u)} \right)^2.$$

We here compare these three estimators on the Normal toy example detailed in the main text with h the projection on both coordinates of the parameter vector θ (see Section 3.1 of the main text). We find that the three estimators behave similarly; boxplots are alike and all tend to overestimate posterior variances (Figure

S1). Results summarized in Table S1 also supports this similarity in NMAE terms. Because the estimator 1 appears to show slightly lower errors for both parameter θ_1 and θ_2 , we decided to use it in the two examples detailed in the main text.

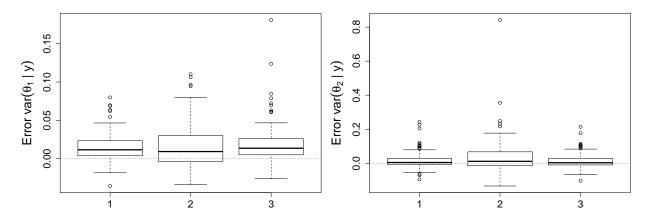


Figure S1: Boxplot comparison of differences between our predictions for $Var(\theta_i \mid y)$ and the true values with the three methods of variance estimation: by reusing weights (method 1, boxplot 1), by building a new RF on square residuals (method 2, boxplot 2) and by using the estimation of the cumulative distribution function (method 3, boxplot 3).

Method	1	2	3
$Var(\theta_1 \mid y)$	0.25	0.28	0.30
$Var(\theta_2 \mid y)$	0.25	0.47	0.25

Table S1: Comparison of normalized mean absolute errors (NMAE) of estimate variances when using three methods, (see legend of Figure S1 and Section 1 of the supplementary material). The smallest NMAE values are in bold characters.

2 Study of covariances of parameters using random forests

2.1 Methodology

We are here interested in another estimate that is frequently produced in a Bayesian analysis, that is the posterior covariance between two univariate transforms of the parameter, $\tau = h(\theta)$ and $\sigma = g(\theta)$ say, $\text{Cov}(\tau, \sigma \mid \eta(y))$. Since we cannot derive this quantity from the approximations to the marginal posteriors of τ and σ , we propose to construct a specific RF for this purpose. With the same notations used in the main text, we denote approximations of posterior expectations for τ and σ , produced by out-of-bag informations, by $\hat{\tau}_{\text{oob}}^{(t)}$ and $\hat{\sigma}_{\text{oob}}^{(t)}$. We use the product of out-of-bag errors for τ and σ in the empirical covariance, and consider $(\tau^{(t)} - \hat{\tau}_{\text{oob}}^{(t)})(\sigma^{(t)} - \hat{\sigma}_{\text{oob}}^{(t)})$ as the response variable. With the previously introduced notations, the corresponding RF estimator is

$$\widetilde{\text{Cov}}(\tau, \sigma \mid \eta(y)) = \frac{1}{B} \sum_{b=1}^{B} \frac{1}{\left| L_{b}(\eta(y)) \right|} \sum_{t: \eta(y^{(t)}) \in L_{b}} n_{b}^{(t)}(\tau^{(t)} - \hat{\tau}_{\text{oob}}^{(t)})(\sigma^{(t)} - \hat{\sigma}_{\text{oob}}^{(t)}).$$

This posterior covariance approximation requires a total of three regression RFs: one for each parameters and one for the covariance approximation.

2.2 Toy regression example

We now apply our RF methodology to a toy regression example for which its non-zero covariance between parameters is the main quantity of interest, hence we consider the case were g and h are the projections on a given coordinate of the parameter vector θ . For a simulated $n \times 2$ design matrix $X = [x_1, x_2]$, we consider the Zellner's hierarchical model (Marin and Robert, 2014, chapter 3)

$$(y_1, \ldots, y_n) \mid \beta_1, \beta_2, \sigma^2 \sim \mathscr{N}_n(X\beta, \sigma^2 Id),$$

 $\beta_1, \beta_2 \mid \sigma^2 \sim \mathscr{N}_2(0, n\sigma^2(X^\top X)^{-1}),$
 $\sigma^2 \sim I\mathscr{G}(4,3),$

where $\mathcal{N}_k(\mu, \Sigma)$ denotes the multivariate normal distribution of dimension k with mean vector μ and covariance matrix Σ , and $I\mathcal{G}(\kappa, \lambda)$ an inverse Gamma distribution with shape parameter κ and scale parameter λ . Provided $X^\top X$ is invertible, this conjugate model leads to closed-form marginal posteriors (Marin and Robert, 2014)

$$\beta_{1}, \beta_{2} \mid y \sim \mathscr{T}_{2} \left(\frac{n}{n+1} (X^{\top}X)^{-1} X^{\top} y, \frac{3+y^{\top} (Id - X(X^{\top}X)^{-1}X^{\top})y/2}{4+n/2} \frac{n}{n+1} (X^{\top}X)^{-1}, 8+n \right),$$

$$\sigma^{2} \mid y \sim I\mathscr{G} \left(4 + \frac{n}{2}, 3 + \frac{1}{2} y^{\top} (Id - X(X^{\top}X)^{-1}X^{\top})y \right),$$

where $\mathscr{T}_k(\mu, \Sigma, \nu)$ is the multivariate Student distribution of dimension k, with location parameter μ , scale matrix Σ and degree of freedom ν .

In our simulation experiment, we concentrate on the non zero covariance of the posterior distribution namely $Cov(\beta_1, \beta_2 \mid y)$. A reference table of N = 10000 replicates of a n-sample with n = 100 is generated. We then create k = 60 summary statistics: the maximum likelihood estimates of β_1 , β_2 , the residual sum of squares, the empirical covariance and correlation between y and x_1 , covariance and correlation between y and

Table S2: Comparison of normalized mean absolute errors (NMAE) of estimate posterior covariances between β_1 and β_2 using random forest (RF), adjusted ridge regression (ARR) and adjusted neural network (ANN) ABC methods. The smallest NMAE value is in bold characters.

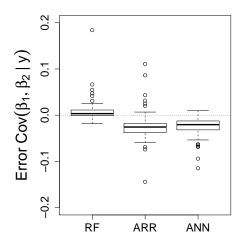


Figure S2: Boxplot comparison of differences between prediction and true values for $Cov(\beta_1, \beta_2 \mid y)$ using random forest (RF), adjusted ridge regression (ARR) and adjusted neural network (ANN) ABC methods.

 x_2 , the sample mean, the sample variance, the sample median, and 50 independent noise variables simulated from a uniform distribution $\mathcal{U}_{[0,1]}$. These noise variables were introduced to be in a sparse context.

Similarly to the Normal example of the main text, we assess the performance of our approach using an independent (Monte Carlo) test dataset of size $N_{\rm pred}=100$ and compare estimation accuracy with the ABC-RF approach from the ones with adjusted ridge regression and neural network ABC methodologies. RF are once again built with B=500 trees, $n_{try}=k/3$ and minimum node size equals to 5 and ABC methods rely on the R package abc with a tolerance parameter equals to 0.1 for ABC methods with adjustment. ABC with neural network adjustment require the specification of the number of layers composing the neural network. We use again 10 layers, the default number of layers in the R package abc. For local linear or ridge regression the corrections are univariate. That is not the case for neural networks which, by construction, perform multivariate correction.

Covariance estimation is a novel feature in this example, Table S2 shows that the ABC-RF approach does better in NMAE terms. As exhibited in Figure S2, ABC-RF overestimates covariances when earlier ABC methods underestimate it. Results are quite encouraging even though we believe the method might still be improved.

3 A basic R code to use the about package version 1.7

We provide some basic R lines of code to use the R package abcrf and conduct RF inference about parameters. There are two possibilities to read simulated data: the user wants to use a *reference table* simulated from the software DIYABC v.2.1.0 (Cornuet *et al.*, 2014) recorded within a .bin file associated with its .txt header file, or the simulated *reference table* is only contained within a .txt file. Of course if the model is simple enough, the user can simulate the *reference table* himself using its own simulator program. In the following, we assume θ is a vector of p parameters and k summary statistics are considered. The # symbol means the text on its right is a comment and ignored by R. We here focus on a single parameter of interest labelled "poi".

Installing and loading the R package aborf

```
install.packages("abcrf") # To install the abcrf package (version 1.7) library(abcrf) # To load the package.
```

Reading data: option 1 - using a .bin and .text files obtained using DIYABC

We assume the *reference table* is recorded within the reftable.bin file and its corresponding header in the header.txt file. The function readRefTable is used to recover the data.

Reading data: option 2 - using a .txt file

We assume that the *reference table* is recorded within yourTxtFile.txt file, composed of a first column corresponding to the scenario indices, p columns of parameters and k columns of summary statistics, the first row is the column labels. The field separator character being a white space.

```
data <- read.table(file = "youTxtFile.txt", header = TRUE, sep = "")
# data is a matrix. The first column is the model indices, the next p are
# the p parameters, the last k are the summary statistics.

index1 <- data[ , 1] == 1 # To store the model 1 indexes.

# We then create a data frame composed of the parameter of interest poi and</pre>
```

Subsetting your dataset

If required, subsetting your datasets stored in data.poi can be easily done with the following line.

```
data.poi <- data.poi[1:10000, ]
# If you are interest in the 10000 first datasets.</pre>
```

Training a random forest

The random forest of the ABC-RF method is built thanks to the regAbcrf function, its principle arguments being a R formula and the corresponding data frame as training dataset. Additional arguments are available, especially the number of trees (ntree, with default values ntree = 500), the minimum node size (min.node.size, with default value min.node.size = 5), and the number of covariates randomly considered at each split (mtry). See the regAbcrf help for further details.

Graphical representations to access the performance of the method

The evolution of the out-of-bag mean squared error depending on the number of tree can be easily represented with the err.regAbcrf function (e.g. Figure 6 of the main text).

The contribution of summary statistics in ABC-RF estimation for the parameter of interest can be retrieved with the plot function applied to an object resulting from regAbcrf.

```
plot(x = model.poi, n.var = 25)
# The contributions of the 25 most important summary statistics are
# represented (e.g. Figure S5).
```

Making predictions

Finally, given a data frame obs.poi containing the summary statistics you want the predictions of posterior quantities of interest for a given dataset (usually the observed dataset). When using DIYABC, note that the summary statistics of the observed dataset are recorded in a file name statobs.txt. The predict method can be used for this purpose. The column names need to be the same than those in the summary statistics of data.poi.

```
# Reading the observed dataset with
obs.poi <- read.table("statobs.txt", skip=2)</pre>
# If obs.poi is not a dataframe or the column names do not match,
# you can use the following lines:
obs.poi <- as.data.frame(obs.poi)</pre>
colnames(obs.poi) <- colnames(data.poi[ ,-1])</pre>
# Prediction is complete by
pred.obsPoi <- predict(object = model.poi, obs = obs.poi,</pre>
                       training = data.poi, quantiles = c(0.025, 0.975),
                        paral = TRUE)
# The 2.5 and 97.5 order quantiles are computed by specifying
# quantiles = c(0.025, 0.975).
# pred.obsPoi is a list containing predictions of interest.
# Posterior mean can be retrieved by
pred.obsPoi$expectation
# Posterior variance by
pred.obsPoi$variance
# Posterior quantiles by
pred.obsPoi$quantiles
```

A graphical representation of the approximate posterior density of poi given obs.poi can be obtained using the densityPlot function.

densityPlot(object = model.poi, obs = obs.poi, training = data.poi, paral = TRUE)

4 Supplementary figures for the Normal toy example

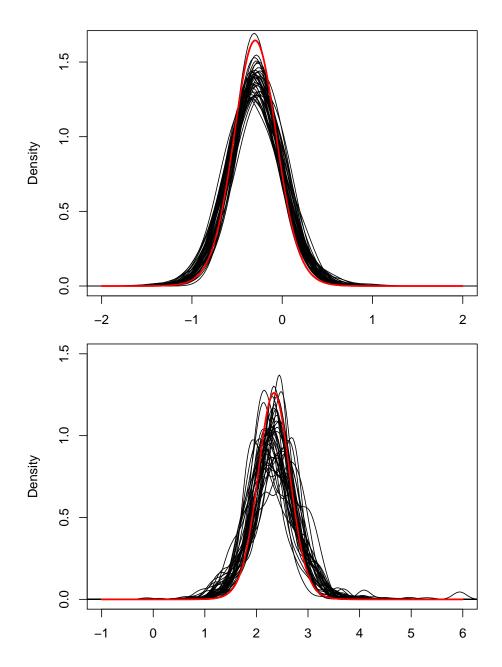


Figure S3: Comparisons of the true posterior density distribution function of θ_1 in the Normal model with a sample of 40 ABC-RF approximations of the posterior density (using RF weights), based on 40 independent *reference tables* and for two different test datasets (upper and lower panels). True posterior densities are represented by red lines.

5 Summary statistics available in the software DIYABC v.2.1.0 for SNP data

For single nucleotide polymorphic (SNP) markers, the program DIYABC v.2.1.0 (Cornuet *et al.*, 2014) proposes a set of summary statistics among those used by population geneticists. These summary statistics are mean values, variance values and proportion of null values across loci, which allow a rough description of the allelic spectrum. Such summary statistics characterize a single, a pair or a trio of population samples.

Single population statistics

HPO_i: proportion of monomorphic loci for population i

HM1_i: mean gene diversity across polymorphic loci (Nei, 1987)

HV1_i: variance of gene diversity across polymorphic loci

HMO_i: mean gene diversity across all loci (Nei, 1987)

Two population statistics

FP0_i&j: proportion of loci with null FST distance between the two samples for populations i and j (Weir and Cockerham, 1984)

FM1_i&j: mean across loci of non null FST distances

FV1_i&j: variance across loci of non null FST distances

FMO_i&j: mean across loci of FST distances (Weir and Cockerham, 1984)

NPO_i&j: proportion of 1 loci with null Nei's distance (Nei, 1972)

NM1_i&j: mean across loci of non null Nei's distances

NV1_i&j: variance across loci of non null Nei's distances

NMO_i&j: mean across loci of Nei's distances (Nei, 1972)

Three population statistics

APO_i_j&k: proportion of loci with null admixture estimate when pop. i comes from an admixture between j and k

AM1_i_j&k: mean across loci of non null admixture estimate

AV1_i_j&k: variance across loci of non null admixture estimated

AMO_i_j&k: mean across all locus admixture estimates (Choisy et al., 2004)

6 Supplementary figures for the Human population genetics example

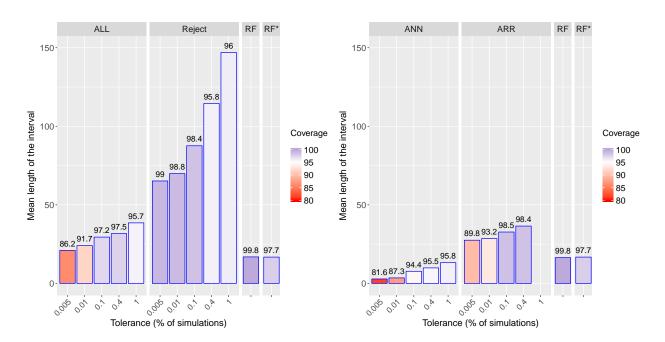


Figure S4: Range and coverage comparison of approximate 95% credible intervals on the ratio N2/Na of the Human population genetics example, obtained with ABC-RF (RF) and with earlier ABC methods: rejection (Reject), adjusted local linear (ALL) or ridge regression (ARR) or neural network (ANN) with various tolerance levels for Reject, ALL, ARR and ANN. Coverages values are specified by bar colors and superimposed values. Heights indicate CI mean lengths. Na is the ancestral African effective population size before the population size change event and N2 the African effective population size after the population size change event (going backward in time). RF* refers to results obtained using ABC-RF when adding 20 additional independent noise variables generated from a uniform $\mathcal{U}_{[0,1]}$ distribution. RF refers to results without noise variables.

7 Contribution of summary statistics in ABC-RF estimation of the parameters ra and N2/Na of the Human population genetics example

In the same spirit than in Pudlo et al. (2016), a by-product of our ABC-RF-based approach is to automatically determine the (most) relevant statistics for the estimation of each parameter by computing a criterion of variable importance (here a variable is a summary statistic). For a summary statistic j, this measure is equal to the total amount of decrease of the residual sum of squares (RSS) due to splits over this given summary statistic, divided by the number of trees.

Indeed, at a given parent node where j is used, with n datasets, and for a given parameter of interest τ , the decrease of the RSS due to the split event is defined by the formula

$$\sum_{i=1}^n (\tau_i - \bar{\tau})^2 - \left(\sum_{i \in \text{ left node}} (\tau_i - \bar{\tau}_L)^2 + \sum_{i \in \text{ right node}} (\tau_i - \bar{\tau}_R)^2\right),$$

where $\bar{\tau}$, $\bar{\tau}_L$ and $\bar{\tau}_R$ are respectively the average of parameter values of the datasets in the parent node, left daughter and right daughter nodes. Computing this decrease among all nodes where a summary statistics j is used, among all the trees of the forest and dividing it by the number of trees is an importance measurement of the covariate j.

Figure S5 shows the contributions of the 30 most important summary statistics (among the 112 statistics proposed by DIYABC) for the ABC-RF estimation of the parameters ra and N2/Na of the Human population genetics example (see Section 3.2 of the main text). The most informative summary statistics are clearly different depending on the parameter of interest. For the admixture rate between two sources populations (ra), all ten most informative statistics correspond to statistics characterizing a pair or a trio of populations (e.g. AV or FMO statistics; see Section 5 of the supplementary materials). Moreover, all those "best" statistics include the populations ASW, GBP and YRI which correspond to the target and the two source populations respectively. On the contrary, for the effective population size ratio N2/Na, seven of the ten most most informative statistics correspond to statistics characterizing within population genetic variation (e.g. HV or HMO; see Section 5 of the supplementary materials). In this case, all those "best" statistics include the African population, which makes sense since N2 is the effective population size in the studied African population and Na in the population ancestral to all studied populations. It is worth stressing that, although the most informative summary statistics make sense in relation to the studied parameters it was difficult if not impossible to a priori and objectively select those statistics. This is not an issue when using the ABC-RF approach as the method automatically extracts the maximum of information from the entire set of proposed statistics.

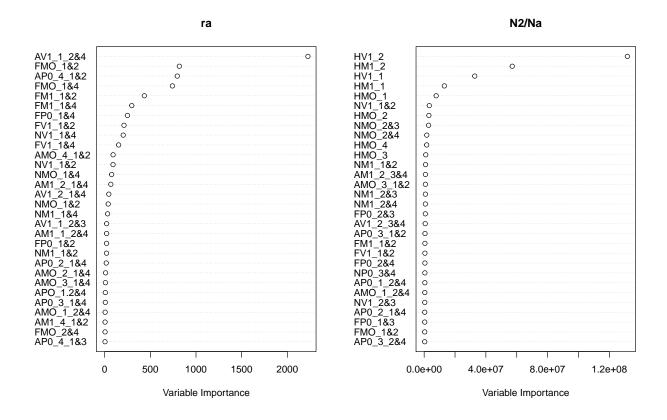


Figure S5: Contributions of the 30 most important summary statistics for the ABC-RF estimation of the parameters ra and N2/Na of the Human population genetics example. The contribution of each statistics is evaluated as the total amount of decrease of the residual sum of squares, divided by the number of trees, for each of the 112 used summary statistics provided for SNP markers by DIYABC. The higher the variable importance the more informative the statistics. The population index(s) is indicated at the end of each statistics. 1 = pop ASW (Americans of African ancestry), 2 = pop YRI (Yoruba, Africa), 3 = pop CHB (Han, Asia) and 4 = pop GBP (British, Europe). For instance FMO_1&4 = mean across loci of Fst distance between the populations 1 and 4. See also Section 5 of the supplementary materials and the Figure 4 of the main text. Note the difference of scale for the importance criterion for parameters ra and N2/Na. This difference can be explained by the difference of scale in parameter values. Indeed, it directly influences the RSS. The parameter ra being bounded in [0,1] contrary to N2/Na, a higher decrease can be expected for the ratio N2/Na than ra.

8 Computation times required by the statistical treatments of the studied methods processed following the generation of the reference table

We here present a comparison of the computation time requirement for the different methods studied in this paper, when predicting estimations of the admixture rate ra in the human population genetics example. ABC methods with rejection or adjusted with local linear regression provide the best results in terms of CPU time even when the tolerance level is equal to 1. The ABC-RF strategy requires moderately higher computing time. The calculation of the RF weights is the most expensive computation part (i.e. 3/4 of computation time). ABC methods using ridge regression or neural network correction become very time consuming when the tolerance level is high.

Method	Tol. level	CPU time (in minutes)
RF	NA	16.64
Reject	0.005	7.54
Reject	0.01	7.54
Reject	0.1	7.78
Reject	0.4	7.98
Reject	1	9.14
ALL	0.005	7.66
ALL	0.01	7.70
ALL	0.1	8.81
ALL	0.4	9.21
ALL	1	11.32
ARR	0.005	6.71
ARR	0.01	6.97
ARR	0.1	40.57
ARR	0.4	560.39
ARR	1	_
ANN	0.005	22.31
ANN	0.01	33.60
ANN	0.1	216.61
ANN	0.4	1160.67
ANN	1	4028.63

Table S3: Comparison of the computation time (in minutes) required - after the gereration of the reference table - for the estimation of the parameter of interest ra on a dataset test table, using ABC-RF (RF), ABC with rejection (Reject), adjusted local linear (ALL), ridge regression (ARR) and neural network (ANN), with various tolerance levels for Reject, ALL, ARR and ANN. The test table included 1000 pseudo-observed datasets and the reference table included 199 000 simulated datasets summarized with 112 statistics. Results were computed on a cluster with 28 CPU cores of 2.4 GHz. NA stands for not appropriate.

References

Choisy, M. et al. (2004). Estimating admixture proportions with microsatellites: comparison of methods based on simulated data. *Mol Ecol*, 13(4):955–968.

Cornuet, J.-M. et al. (2014). DIYABC v2.0: a software to make approximate Bayesian computation inferences about population history using single nucleotide polymorphism, DNA sequence and microsatellite data. *Bioinformatics*, 30(8):1187–1189.

Marin, J.-M. and Robert, C. P. (2014) Bayesian Essentials with R. Springer.

Nei, M. (1972). Genetic distance between populations. *The American Naturalist*, 106(949):283–292.

Nei, M. (1987). Molecular Evolutionary Genetics. Columbia University Press, New York, USA, first edition.

Pudlo, P. et al. (2016) Reliable ABC model choice via random forests. *Bioinformatics*, 32(6):859-866.

Weir, B. S. and Cockerham, C. C. (1984). Estimating F-Statistics for the Analysis of Population Structure. *Evolution*, 38(6):1358–1370.